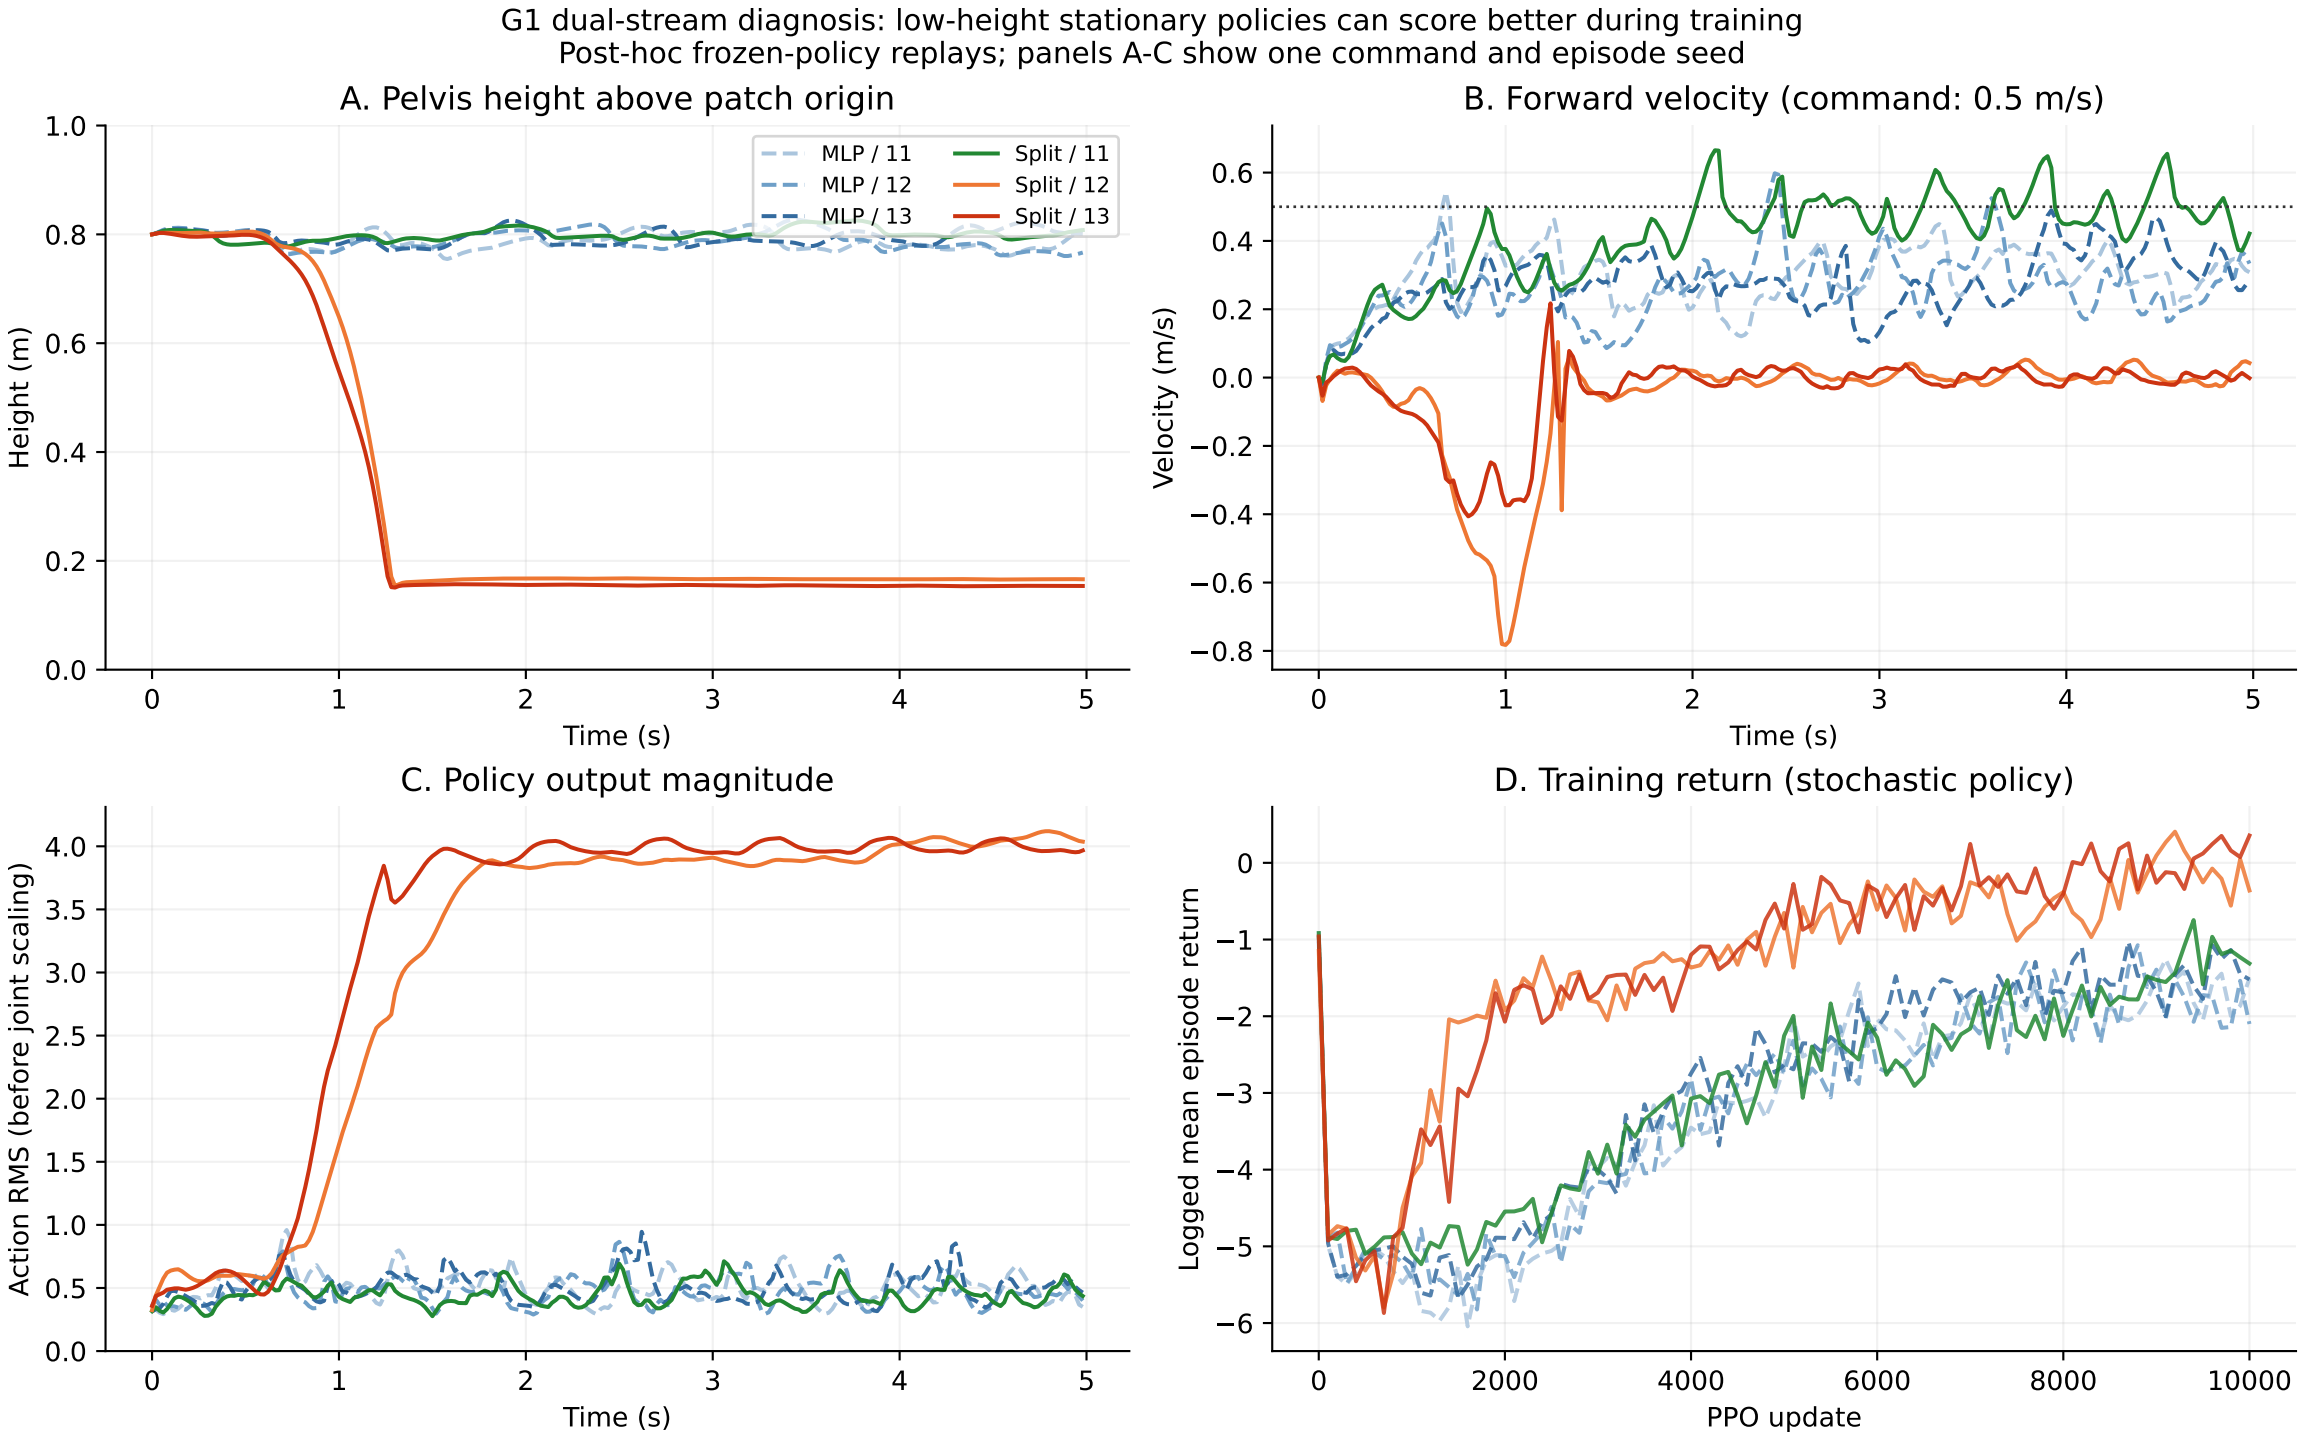
G1 dual-stream diagnosis: low-height stationary policies can score better during training
Post-hoc frozen-policy replays; panels A-C show one command and episode seed
A. Pelvis height above patch origin
Height (m)
Time (s)
MLP / 11
MLP / 12
MLP / 13
Split / 11
Split / 12
Split / 13
B. Forward velocity (command: 0.5 m/s)
Velocity (m/s)
Time (s)
C. Policy output magnitude
Action RMS (before joint scaling)
Time (s)
D. Training return (stochastic policy)
Logged mean episode return
PPO update